



Indian Institute of Technology Delhi

Operating System

ELL 783

Assignment 2

Name: Ranjan Kumar Singha
Entry No: 2023EET2192

Name: Atul Kumar Rana
Entry No: 2023EET2195

Course Instructor: Dr. Smruti R. Sarangi

1 Additional Process Attributes

Added the following fields to `proc.h` to keep track of various attributes needed to implement the required scheduling algorithms and system calls.

- `int sched_policy`:: Scheduling policy of the process (-1: XV6 default policy or 0: EDF or 1: RMS or 4: non-schedulable processes). It is set by a user process using the `sched_policy(pid, value)` system call.
- `int elapsed_time`:: Elapsed time of the process. It counts the number of ticks for which the process was in `RUNNING` state.
- `int execution_time`:: Total allowed execution time of the process. It is set by a user process using the `execution_time(pid, value)` system call.
- `int deadline`:: Hard Deadline of the process. Any new process that fails the schedulability check is killed so that no deadlines for accepted processes are sacrificed. It is set using the `deadline(pid, value)` system call.
- `int rate`:: Rate of the assumed periodic processes. It is set by a user process using the `rate(pid, value)` system call.
- `int priority`:: Current priority level of each process (1-3) (higher value represents lower priority). It is calculated using the rate of the process.
- `unsigned int arrival_time`:: Start time of the process. It records the tick value when the process's `sched_policy` was set.

```
1  int sched_policy;           // Schedule policy
2  int elapsed_time;          // elapsed time
3  int execution_time;        // execution time
4  int deadline;              // deadline
5  int rate;                  // rate
6  int priority;              // priority
7  unsigned int arrival_time; // arrival time
```

Code Snippet in `proc.c`

Initialized the count of EDF and RMS

```
1  int edf_count=0; // For how many process I have to run EDF
2  int rms_count=0; // For how many process I have to run RMS
```

This code snippet is responsible for decrementing the counts of processes remaining to be executed under the EDF and RMS scheduling algorithms after the completion of a process execution.

```

1 // After complete execution of a process, decrement the time
  count of the algorithm that has to be implemented.
2 if(curproc->sched_policy==0){
3     edf_count = edf_count - 1;
4 }
5 else if(curproc->sched_policy==1){
6     rms_count = rms_count - 1;
7 }

```

In the `proc.c` file, the `userinit()` and `fork()` functions initialize critical attributes of the process structure (`proc`) to ensure proper initialization for every process, whether created during system boot or through forking.

```

1 void userinit(void) {
2     // Other initialization code...
3
4     // Necessary variables needed for Parent Process
5     p->sched_policy = -1;
6     p->execution_time = -1;
7     p->deadline = -1;
8     p->elapsed_time = 0;
9
10    // Other initialization code...
11 }

```

```

1 int fork(void) {
2     // Other initialization code...
3
4     // Necessary variables for the newly created Child
      processes
5     np->sched_policy = -1;
6     np->execution_time = -1;
7     np->elapsed_time = 0;
8
9     // Other initialization code...
10 }

```

Code Snippet in `trap.c`

```

1 if (myproc() && myproc()->state == RUNNING &&
2 tf->trapno == T_IRQ0 + IRQ_TIMER)
3 {

```

```

4     if ((myproc()->sched_policy >= 0) &&
5         (myproc()->elapsed_time >= myproc()->exec_time))
6     {
7         cprintf("The arrival time and pid value of the
8             completed process is %d %d \n", myproc()->
9             arrival_time, myproc()->pid);
10        exit();
11    }
12    else
13        yield();
14}

```

In the modified `trap.c` file, the above code snippet is used to handle processes. If the current process (`myproc()`) is in the `RUNNING` state and a timer interrupt occurs, the code checks if the process has become unwanted or completed. If the process has become unwanted (i.e., its `sched_policy` is not equal to `-1` and its `elapsed_time` exceeds its `exec_time`), the process is terminated and its details (arrival time and PID) are printed. Otherwise, the process yields the CPU.

2 sys_sched_policy

EDF Code Snippet

```

1  if (policy == 0) { // EDF
2      int edf_check = handle_edf_policy(pid);
3      if (edf_check != 0) {
4          // Terminate process if EDF policy check fails
5          for (p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
6              if (p->pid == pid) {
7                  p->state = ZOMBIE;
8                  break;
9              }
10         }
11         release(&ptable.lock);
12         return edf_check; // Return error code
13     }
14 }

```

- The code snippet checks if the specified scheduling policy is EDF (`policy == 0`).
- It calls the `handle_edf_policy` function to perform policy-specific checks.
- The `handle_edf_policy` function calculates the CPU utilization and checks if it exceeds 100

- If the check fails (indicating a violation of scheduling constraints), the process with the specified PID is terminated by setting its state to ZOMBIE.

RMS Code Snippet

```

1 else if (policy == 1) { // RMS
2     // Count RMS processes
3     int count = 0;
4     for (p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
5         if (p->sched_policy == 1 || p->pid == pid) {
6             count++;
7         }
8     }
9     // Check RMS policy constraints
10    int rms_check = handle_rms_policy(pid, count);
11    if (rms_check != 0) {
12        // Terminate process if RMS policy check fails
13        for (p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
14            if (p->pid == pid) {
15                p->state = ZOMBIE;
16                break;
17            }
18        }
19        release(&ptable.lock);
20        return rms_check;
21    }
22 }

```

- The code snippet checks if the specified scheduling policy is RMS (policy == 1).
- It counts the number of processes currently using the RMS policy or having the specified process ID.
- It calculates the total resource consumed by RMS processes and checks if it exceeds the threshold value.
- If the check fails (indicating a violation of resource limits), the process with the specified PID is terminated by setting its state to ZOMBIE.

handle_edf_policy Function

The `handle_edf_policy` function calculates the CPU utilization for processes adhering to the Earliest Deadline First (EDF) policy.

- CPU Utilization Calculation:
 - It iterates through all processes in the process table.

- For each process with the EDF policy or the specified process ID, it calculates CPU utilization as the ratio of execution time to deadline.
- Violation Check:
 - If the total CPU utilization exceeds 100

```

1 int handle_edf_policy(int pid) {
2     float cpu_utilization = 0.0;
3     struct proc *p;
4     for (p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
5         if ((p->sched_policy == 0 && p->state != UNUSED) || p->
6             pid == pid) {
7             cpu_utilization += (float)(p->execution_time) / (
8                 float)(p->deadline);
9         }
10    }
11    if (cpu_utilization > 1.0) {
12        return -22; // Return error code indicating Not
13        Schedulable
14    }
15    return 0;
16 }

```

handle_rms_policy Function

The `handle_rms_policy` function calculates the total resource consumed by processes adhering to the Rate-Monotonic Scheduling (RMS) policy.

- Resource Calculation:
 - It iterates through all processes in the process table.
 - For each process with the RMS policy or the specified process ID, it calculates resource consumption as the product of execution time and rate.
- Resource Limit Check:
 - If the total resource consumption exceeds the threshold value determined by the rate-monotonic bound, the function returns an error code (-22) to indicate that the resource limit has been exceeded.

```

1 int handle_rms_policy(int pid, int count) {
2     float total_resource = 0.0;
3     struct proc *p;
4     for (p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
5         if (p->sched_policy == 1 || p->pid == pid) {

```

```

6         total_resource += ((float)(p->execution_time) * (
           float)(p->rate));
7     }
8 }
9 if (total_resource > (threshold_array[count - 1] * 100.0))
10 {
11     return -22; // Return error code indicating resource
           limit exceeded
12 }
13 return 0;
14 }

```

RMS Threshold Array

The `threshold_array` stores Luiland bound values used in Rate-Monotonic Scheduling (RMS). Each element of the array corresponds to a specific count of processes, representing the threshold value beyond which the system may become overloaded.

- ****Threshold Values****:
 - The `threshold_array` contains predefined threshold values based on Luiland's bound.
 - Each element in the array corresponds to a count of processes, starting from 1.
 - These values are derived from the rate-monotonic scheduling theory and are used to determine the maximum total resource consumption allowed for a given number of processes.

RMS Threshold Array

```

1 // Array to store Luiland bound values
2 float threshold_array[64] = {
3     1.0000, 0.8284, 0.7798, 0.7568, 0.7435, 0.7348, 0.7286,
4     0.7241,
5     ...,
6     0.6974, 0.6973, 0.6972, 0.6972, 0.6971, 0.6970, 0.6970,
7     0.6969
8 };

```

3 sys_exec_time Function

- This function sets the execution time for a specific process identified by its PID.
- It takes two arguments: the PID of the process and the execution time to be set.

- If successful, it updates the execution time of the process and returns 0. If unsuccessful, it returns -1.

```

1 int
2 sys_exec_time(void) {
3     int pid, exec_time;
4     if(argint(0, &pid) < 0 || argint(1, &exec_time) < 0){
5         return -1;
6     }
7     else{
8         struct proc *p;
9         acquire(&ptable.lock);
10        for (p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
11            if (p->pid == pid) {
12                p->execution_time = exec_time;
13                release(&ptable.lock);
14                return 0;
15            }
16        }
17        release(&ptable.lock);
18        return -1;
19    }
20 }

```

4 sys_deadline Function

- This function sets the deadline for a specific process identified by its PID.
- It takes two arguments: the PID of the process and the deadline to be set.
- If successful, it updates the deadline of the process and returns 0. If unsuccessful, it returns -1.

```

1 int
2 sys_deadline(void){
3     int pid, deadline;
4     if(argint(0, &pid) < 0 || argint(1, &deadline) < 0){
5         return -1;
6     }
7     else{
8         struct proc *p;
9         acquire(&ptable.lock);
10        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
11            if(p->pid == pid){
12                p->deadline = deadline;

```



```

13         release(&ptable.lock);
14         return 0;
15     }
16 }
17 release(&ptable.lock);
18 return -1;
19 }
20 }

```

5 sys_rate Function

- This function sets the rate for a specific process identified by its PID.
- It calculates the priority of the process based on the rate and sets it accordingly.
- If successful, it updates the rate and priority of the process and returns 0. If unsuccessful, it returns -1.

```

1 int
2 sys_rate(void){
3     int pid,rate;
4     if(argint(0, &pid) < 0 || argint(1, &rate) < 0){
5         return -1;
6     }
7     else{
8         struct proc *p;
9         acquire(&ptable.lock);
10        for(p=ptable.proc;p<&ptable.proc[NPROC];p++){
11            if(p->pid==pid){
12                p->rate=rate;
13                p->priority=max_value(1,ceil_value(((30.0-rate)
14                    /29.0)*3.0));
15                release(&ptable.lock);
16                return 0;
17            }
18        }
19        release(&ptable.lock);
20        return -1;
21    }
22 }

```

Formula for the weight (w) for a process with rate (r):

$$w = \max\left(1, \left\lceil \frac{30 - \text{rate}}{29} \times 3 \right\rceil\right)$$

```

1 // Functions for necessary operations
2 int
3 max_value(int num1,int num2){
4     return num1>num2 ? num1 : num2;
5 }
6
7 int
8 ceil_value(float num){
9     return (int)num+1;
10 }

```

6 Screenshot

